

# Project\_College\_Admission.R

might

2022-10-30

```
#importing library  
library(ggplot2)  
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(caTools)  
library(ROCR)  
library(MASS)  
library(caret)
```

```
## Loading required package: lattice
```

```
library(rpart)  
library(rpart.plot)  
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```
library(kernlab)
```

```
##  
## Attaching package: 'kernlab'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## alpha
```

```
library(readr)
library(ggpubr)

#Clearing the r environment
rm(list = ls(all = TRUE))

#loading data
clg_ad <- read.csv("College_admission.csv")
head(clg_ad)
```

```
##   admit gre  gpa ses Gender_Male Race rank
## 1     0 380 3.61  1           0    3    3
## 2     1 660 3.67  2           0    2    3
## 3     1 800 4.00  2           0    2    1
## 4     1 640 3.19  1           1    2    4
## 5     0 520 2.93  3           1    2    4
## 6     1 760 3.00  2           1    1    2
```

```
tail(clg_ad)
```

```
##   admit gre  gpa ses Gender_Male Race rank
## 395     1 460 3.99  3           1    3    3
## 396     0 620 4.00  2           0    2    2
## 397     0 560 3.04  2           0    1    3
## 398     0 460 2.63  3           0    2    2
## 399     0 700 3.65  1           1    1    2
## 400     0 600 3.89  2           1    3    3
```

```
#descriptive statistics
summary(clg_ad$gpa)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  2.260   3.130   3.395   3.390   3.670   4.000
```

```
#1. Find the missing values. (if any, perform missing value treatment)
```


```
sum(is.null(clg_ad))
```

```
## [1] 0
```

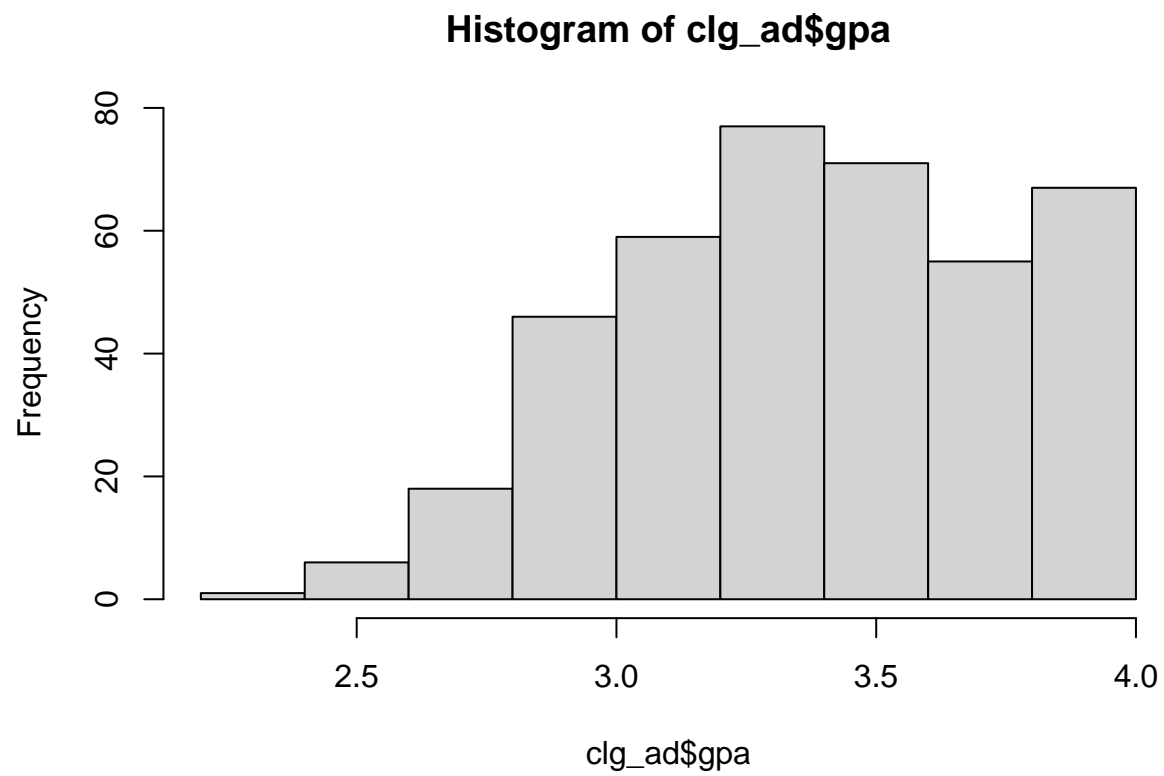
```
#Since there are no Null values in the dataset, no need for missing value treatment
```

```
#2. Find outliers (if any, then perform outlier treatment).
#Visualizing continuous variables
```


```
hist(clg_ad$gre)
```




```
hist(clg_ad$gpa)
```



```
#Using Boxplot to understand if there are any outliers present in GRE variable  
boxplot(clg_ad$gre, horizontal = T)
```



```
# By looking at the Boxplot, we can see there are outliers in GRE variable  
greoutlier <- boxplot(clg_ad$gre)$out
```



```
#checking length of GRE outliers
length(greoutlier)
```

```
## [1] 4
```

```
#Removing outliers from GRE variable
```

```
Q <- quantile(clg_ad$gre, probs=c(.25, .75), na.rm = FALSE)
iqr <- IQR(clg_ad$gre)
```


```
up <- Q[2]+1.5*iqr # Upper Range
```

```
low<- Q[1]-1.5*iqr # Lower Range
```


```
clg_ad <- subset(clg_ad, clg_ad$gre > (Q[1] - 1.5*iqr) & clg_ad$gre < (Q[2]+1.5*iqr))
head(clg_ad)
```

```
##   admit gre  gpa ses Gender_Male Race rank
## 1     0 380 3.61  1           0    3    3
## 2     1 660 3.67  2           0    2    3
## 3     1 800 4.00  2           0    2    1
## 4     1 640 3.19  1           1    2    4
## 5     0 520 2.93  3           1    2    4
## 6     1 760 3.00  2           1    1    2
```

```
#Using Boxplot to understand if there are any outliers present in GPA variable  
boxplot(clg_ad$gpa, horizontal = T)
```



```
# By looking at the Boxplot, we can see there are outliers in GPA variable  
gpaoutlier <- boxplot(clg_ad$gpa)$out
```



```
length(gpaoutlier)
```

```
## [1] 1
```

```
#There is only one outlier present in this Variable, we can remove it.
```

```
#Removing outliers from GPA variable
```

```
Q <- quantile(clg_ad$gpa, probs=c(.25, .75), na.rm = FALSE)
```

```
iqr <- IQR(clg_ad$gpa)
```

```
up <- Q[2]+1.5*iqr # Upper Range
```

```
low<- Q[1]-1.5*iqr # Lower Range
```

```
clg_ad <- subset(clg_ad, clg_ad$gpa > (Q[1] - 1.5*iqr) & clg_ad$gpa < (Q[2]+1.5*iqr))  
head(clg_ad)
```

```
##   admit gre  gpa ses Gender_Male Race rank  
## 1     0 380 3.61  1           0    3    3  
## 2     1 660 3.67  2           0    2    3  
## 3     1 800 4.00  2           0    2    1  
## 4     1 640 3.19  1           1    2    4  
## 5     0 520 2.93  3           1    2    4  
## 6     1 760 3.00  2           1    1    2
```



*#now outliers have been removed from data, After removing outliers we have 395 data points*

*#3. Find the structure of the data set and if required, transform the numeric data type to factor and v*  
*#structure of data*

```
str(clg_ad)
```

```
## 'data.frame':    395 obs. of  7 variables:
## $ admit      : int  0 1 1 1 0 1 1 0 1 0 ...
## $ gre        : int 380 660 800 640 520 760 560 400 540 700 ...
## $ gpa        : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
## $ ses        : int  1 2 2 1 3 2 2 2 1 1 ...
## $ Gender_Male: int  0 0 0 1 1 1 1 0 1 0 ...
## $ Race       : int  3 2 2 2 2 1 2 2 1 2 ...
## $ rank       : int  3 3 1 4 4 2 1 2 3 2 ...
```


```
clg_ad$admit <- as.factor(clg_ad$admit)
clg_ad$ses <- as.factor(clg_ad$ses)
clg_ad$Gender_Male <- as.factor(clg_ad$Gender_Male)
clg_ad$Race <- as.factor(clg_ad$Race)
clg_ad$rank <- as.factor(clg_ad$rank)
str(clg_ad)
```

```
## 'data.frame':    395 obs. of  7 variables:
## $ admit      : Factor w/ 2 levels "0","1": 1 2 2 2 1 2 2 1 2 1 ...
## $ gre        : int 380 660 800 640 520 760 560 400 540 700 ...
## $ gpa        : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
## $ ses        : Factor w/ 3 levels "1","2","3": 1 2 2 1 3 2 2 2 1 1 ...
## $ Gender_Male: Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 2 1 ...
## $ Race       : Factor w/ 3 levels "1","2","3": 3 2 2 2 2 1 2 2 1 2 ...
## $ rank       : Factor w/ 4 levels "1","2","3","4": 3 3 1 4 4 2 1 2 3 2 ...
```


*#4. Find whether the data is normally distributed or not. Use the plot to determine the same.*

*#checking distribution of GRE variable*

```
plot(density(clg_ad$gre))
```



```
ggqqplot(clg_ad$gre)
```




*#Since Majority of the data points fall on the line,  
#we can assume that the data is normally distributed*


*##Normality test using Shapiro test*  
shapiro.test(clg\_ad\$gre)

```
##
##  Shapiro-Wilk normality test
##
## data:  clg_ad$gre
## W = 0.98282, p-value = 0.0001223
```

*#checking distribution of GPA variable*  
plot(density(clg\_ad\$gpa))



```
##Since Majority of the data points fall on the line, we can assume that the data is normally distributed  
ggqqplot(clg_ad$gpa)
```



```
#Normality test using Shapiro test
shapiro.test(clg_ad$gpa)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  clg_ad$gpa
## W = 0.97646, p-value = 5.004e-06
```

```
#5. Normalize the data if not normally distributed
#Since the data is normally distributed, data transformation is not required
#But however there is variance in GRE and GPA variable, we can clearly see
#that gre values are in hundred times more than the gpa values.
#In this case when we build model, gre will internally influence
#the result more due to its larger value.
#To avoid problem and for accurate model,
#we can scale down the variables to avoid problem and accurate model.
```

```
#Creating a copy of the data
clgad <- clg_ad
clg_ad$gre <- scale(clg_ad$gre, center = T, scale = T)
clg_ad$gpa <- scale(clg_ad$gpa, center = T, scale = T)
head(clg_ad)
```

```
##   admit      gre      gpa ses Gender_Male Race rank
```

```
## 1      0 -1.8927212  0.5655683  1          0   3   3
## 2      1  0.6160871  0.7254249  2          0   2   3
## 3      1  1.8704913  1.6046358  2          0   2   1
## 4      1  0.4368865 -0.5534273  1          1   2   4
## 5      0 -0.6383171 -1.2461390  3          1   2   4
## 6      1  1.5120901 -1.0596397  2          1   1   2
```

*#6. Use variable reduction techniques to identify significant variables.*

*#We can build logistic regression model and identify significant variables.*

```
set.seed(1234)
sampledata <- sample.split(clg_ad$admit, SplitRatio = 0.7)
train <- clg_ad[sampledata==T,]
test <- clg_ad[sampledata==F,]

test_without_admit <- test[,-1]

# fit the model
class(train$admit)
```

```
## [1] "factor"
```

```
log_reg <- glm(admit ~ . , data = train, family = 'binomial')
summary(log_reg)
```

```
##
## Call:
## glm(formula = admit ~ ., family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8138  -0.8617  -0.5745   1.0071   2.2491
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.7374    0.4309   1.711 0.087000 .
## gre             0.3871    0.1556   2.488 0.012833 *
## gpa             0.2469    0.1525   1.619 0.105418
## ses2           -0.2934    0.3504  -0.837 0.402436
## ses3           -0.2998    0.3430  -0.874 0.382196
## Gender_Male1  -0.2494    0.2898  -0.861 0.389398
## Race2         -0.5247    0.3472  -1.511 0.130761
## Race3         -0.3224    0.3429  -0.940 0.347027
## rank2         -0.7151    0.3855  -1.855 0.063561 .
## rank3         -1.5240    0.4226  -3.606 0.000311 ***
## rank4         -2.0875    0.5807  -3.595 0.000325 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 345.55  on 275  degrees of freedom
## Residual deviance: 301.79  on 265  degrees of freedom
```

```
## AIC: 323.79
##
## Number of Fisher Scoring iterations: 4
```

```
#Using step AIC method to identify significant variables
model_AIC = stepAIC(object = log_reg,direction = "both")
```

```
## Start: AIC=323.79
## admit ~ gre + gpa + ses + Gender_Male + Race + rank
##
##           Df Deviance    AIC
## - ses      2   302.78 320.78
## - Race     2   304.16 322.16
## - Gender_Male 1   302.53 322.53
## <none>      3   301.79 323.79
## - gpa      1   304.44 324.44
## - gre      1   308.24 328.24
## - rank     3   323.76 339.76
##
## Step: AIC=320.78
## admit ~ gre + gpa + Gender_Male + Race + rank
##
##           Df Deviance    AIC
## - Gender_Male 1   303.44 319.44
## - Race       2   305.47 319.47
## <none>        3   302.78 320.78
## - gpa        1   305.40 321.40
## + ses        2   301.79 323.79
## - gre        1   309.18 325.18
## - rank       3   325.26 337.26
##
## Step: AIC=319.44
## admit ~ gre + gpa + Race + rank
##
##           Df Deviance    AIC
## - Race       2   305.88 317.88
## <none>        3   303.44 319.44
## - gpa        1   306.15 320.15
## + Gender_Male 1   302.78 320.78
## + ses        2   302.53 322.53
## - gre        1   309.63 323.63
## - rank       3   325.48 335.48
##
## Step: AIC=317.88
## admit ~ gre + gpa + rank
##
##           Df Deviance    AIC
## <none>        3   305.88 317.88
## - gpa        1   308.62 318.62
## + Race       2   303.44 319.44
## + Gender_Male 1   305.47 319.47
## + ses        2   304.69 320.69
## - gre        1   311.84 321.84
## - rank       3   328.35 334.35
```

```
summary(model_AIC)
```

```
##
## Call:
## glm(formula = admit ~ gre + gpa + rank, family = "binomial",
##      data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6582  -0.8584  -0.5817   1.0534   2.3413
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.1806     0.3145   0.574 0.565843
## gre           0.3652     0.1524   2.396 0.016563 *
## gpa           0.2482     0.1511   1.642 0.100525
## rank2        -0.7396     0.3751  -1.972 0.048646 *
## rank3        -1.5088     0.4161  -3.626 0.000287 ***
## rank4        -2.1181     0.5753  -3.682 0.000231 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 345.55  on 275  degrees of freedom
## Residual deviance: 305.88  on 270  degrees of freedom
## AIC: 317.88
##
## Number of Fisher Scoring iterations: 4
```

```
#From the results above GRE, GPA and Rank are the most significant variables.
```

```
#7. Run logistic model to determine the factors that
#influence the admission process of a student (Drop insignificant variables)
summary(log_reg)
```

```
##
## Call:
## glm(formula = admit ~ ., family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8138  -0.8617  -0.5745   1.0071   2.2491
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.7374     0.4309   1.711 0.087000 .
## gre           0.3871     0.1556   2.488 0.012833 *
## gpa           0.2469     0.1525   1.619 0.105418
## ses2        -0.2934     0.3504  -0.837 0.402436
## ses3        -0.2998     0.3430  -0.874 0.382196
```



```
## Gender_Male1 -0.2494      0.2898 -0.861 0.389398
## Race2        -0.5247      0.3472 -1.511 0.130761
## Race3        -0.3224      0.3429 -0.940 0.347027
## rank2        -0.7151      0.3855 -1.855 0.063561 .
## rank3        -1.5240      0.4226 -3.606 0.000311 ***
## rank4        -2.0875      0.5807 -3.595 0.000325 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 345.55  on 275  degrees of freedom
## Residual deviance: 301.79  on 265  degrees of freedom
## AIC: 323.79
##
## Number of Fisher Scoring iterations: 4
```

*#After looking at the summary results obtained from logistic regression,  
#factors that influence admission processs are GRE, GPA and RANK*

```
# Dropping variables that are insignificant.
clg_ad1 <- subset(clg_ad, select = -c(ses,Gender_Male,Race))
head(clg_ad1)
```

```
##   admit      gre      gpa rank
## 1     0 -1.8927212  0.5655683   3
## 2     1  0.6160871  0.7254249   3
## 3     1  1.8704913  1.6046358   1
## 4     1  0.4368865 -0.5534273   4
## 5     0 -0.6383171 -1.2461390   4
## 6     1  1.5120901 -1.0596397   2
```

*#8. Calculate the accuracy of the model and run validation techniques  
#Using the result of logistic regresssion model to calculate the accuracy  
#Using only the significant variables for Building models*

```
class(clg_ad1)
```

```
## [1] "data.frame"
```

```
head(clg_ad1)
```

```
##   admit      gre      gpa rank
## 1     0 -1.8927212  0.5655683   3
## 2     1  0.6160871  0.7254249   3
## 3     1  1.8704913  1.6046358   1
## 4     1  0.4368865 -0.5534273   4
## 5     0 -0.6383171 -1.2461390   4
## 6     1  1.5120901 -1.0596397   2
```

```
set.seed(123)
```

```

sample_data <- sample.split(clg_ad1$admit, SplitRatio = 0.7)
Train <- clg_ad1[sample_data==T,]
Test <- clg_ad1[sample_data==F,]

test_without_admit <- Test[,-1]

logreg_model <- glm(admit ~ . , data = Train, family = 'binomial')

summary(logreg_model)

```

```

##
## Call:
## glm(formula = admit ~ ., family = "binomial", data = Train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.4134  -0.8949  -0.6772   1.2222   2.0431
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.2022     0.3247  -0.623   0.5334
## gre           0.2153     0.1411   1.526   0.1269
## gpa           0.2224     0.1421   1.565   0.1177
## rank2        -0.2949     0.3792  -0.778   0.4368
## rank3        -1.0567     0.4252  -2.485   0.0130 *
## rank4        -1.2841     0.5104  -2.516   0.0119 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 345.55  on 275  degrees of freedom
## Residual deviance: 326.50  on 270  degrees of freedom
## AIC: 338.5
##
## Number of Fisher Scoring iterations: 4

```

```

#Predicting on test data and calculating the accuracy using confusion matrix
## Model Evaluation
prob_train <- predict(logreg_model,newdata = Test[,-1], type = "response")
preds_train <- ifelse(prob_train > 0.49,1,0) # use 0.5 or 0.49 to get the best accuracy
comp = table(Test$admit,preds_train)
confusionMatrix(comp,positive = "0")

```

```

## Confusion Matrix and Statistics
##
##      preds_train
##      0  1
## 0  77  4
## 1  28 10
##
##
##              Accuracy : 0.7311

```

```
##          95% CI : (0.6421, 0.8082)
##    No Information Rate : 0.8824
##    P-Value [Acc > NIR] : 1
##
##          Kappa : 0.2568
##
##    McNemar's Test P-Value : 4.785e-05
##
##          Sensitivity : 0.7333
##          Specificity : 0.7143
##          Pos Pred Value : 0.9506
##          Neg Pred Value : 0.2632
##          Prevalence : 0.8824
##          Detection Rate : 0.6471
##    Detection Prevalence : 0.6807
##          Balanced Accuracy : 0.7238
##
##          'Positive' Class : 0
##
```

*#From the above model we can see that  
#logistic regression model is predicting the admission rate is 73.11%*

*#9. Try other modelling techniques like decision tree and SVM and select a champion model*

*#Decision Tree*  
*#Using library C50*  
library(C50)

*#Building the model and printing the summary*  
c5\_tree\_model <- C5.0(admit~., Train, rules = T)  
c5\_tree\_model

```
##
## Call:
## C5.0.formula(formula = admit ~ ., data = Train, rules = T)
##
## Rule-Based Model
## Number of samples: 276
## Number of predictors: 3
##
## Number of Rules: 0
##
## Non-standard options: attempt to group attributes
```

```
summary(c5_tree_model)
```

```
##
## Call:
## C5.0.formula(formula = admit ~ ., data = Train, rules = T)
```

```
##
##
## C5.0 [Release 2.07 GPL Edition]      Sun Oct 30 21:08:29 2022
## -----
##
## Class specified by attribute 'outcome'
##
## Read 276 cases (4 attributes) from undefined.data
##
## Rules:
##
## Default class: 0
##
##
## Evaluation on training data (276 cases):
##
##      Rules
##      -----
##      No      Errors
##
##      0      88(31.9%)    <<
##
##      (a)   (b)    <-classified as
##      ----  ----
##      188          (a): class 0
##      88           (b): class 1
##
##
## Time: 0.0 secs
```

#### *#Predicting on test data*

```
prob_pred <- predict(c5_tree_model, Test[, -1])
prob_pred
```

```
##      [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##     [38] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##     [75] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##    [112] 0 0 0 0 0 0 0 0
## Levels: 0 1
```

#### *#Using confusion metric to calculate accuracy*

```
confusionMatrix(prob_pred, Test$admit)
```

#### ## Confusion Matrix and Statistics

```
##
##      Reference
## Prediction  0  1
##           0 81 38
##           1  0  0
##
##      Accuracy : 0.6807
##      95% CI : (0.589, 0.7631)
```

```
##      No Information Rate : 0.6807
##      P-Value [Acc > NIR] : 0.5438
##
##              Kappa : 0
##
##  Mcnemar's Test P-Value : 1.947e-09
##
##      Sensitivity : 1.0000
##      Specificity : 0.0000
##      Pos Pred Value : 0.6807
##      Neg Pred Value :      NaN
##      Prevalence : 0.6807
##      Detection Rate : 0.6807
##      Detection Prevalence : 1.0000
##      Balanced Accuracy : 0.5000
##
##      'Positive' Class : 0
##
```

*#using Decision Tree (library C50) is giving 68.07% accuracy*

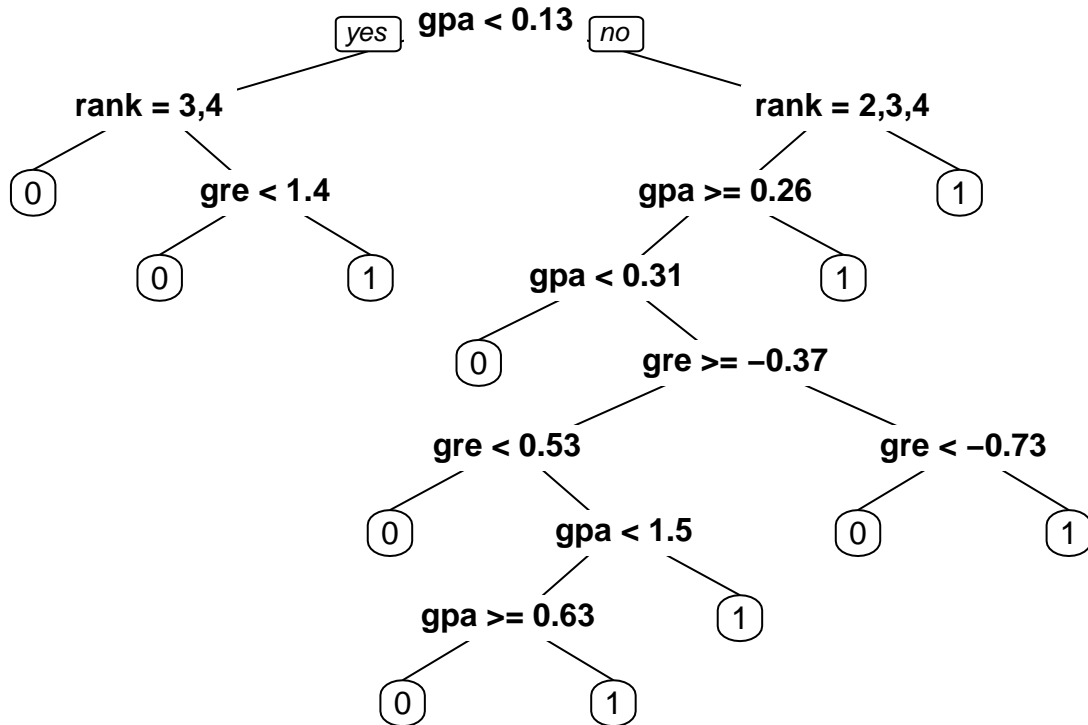
*#Decision Tree using rpart*

*#Model Building*

```
rpart_tree_model <- rpart(admit ~ .,
                          data = Train,
                          method = "class")
```

*# display decision tree*

```
prp(rpart_tree_model)
```



```
# make predictions on the test set
```

```
tree_predict <- predict(rpart_tree_model, Test[,-1], type = "class")
tree_predict
```

```
## 3 5 9 10 11 13 15 16 19 23 26 31 36 37 38 44 48 53 55 57
## 1 0 0 0 1 1 1 0 0 0 1 1 0 0 0 0 0 0 0 0
## 58 61 62 64 77 78 80 82 88 89 92 95 98 100 101 102 104 106 107 109
## 0 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0 1 0 1 0
## 124 128 129 130 139 148 152 154 156 158 160 164 165 168 176 178 179 185 187 189
## 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0
## 193 194 196 203 207 208 213 214 218 220 239 242 247 248 256 260 262 263 271 275
## 0 0 0 1 1 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0
## 277 279 283 286 288 297 301 303 314 315 321 324 325 327 329 330 337 338 340 348
## 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0
## 350 353 355 361 363 365 366 367 368 372 374 379 384 386 388 390 391 393 394
## 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0
## Levels: 0 1
```

```
# evaluate the results
```

```
confusionMatrix(tree_predict, as.factor(Test$admit), positive = "0")
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction 0 1
```

```
##          0 72 20
##          1  9 18
##
##          Accuracy : 0.7563
##          95% CI : (0.6691, 0.8303)
##    No Information Rate : 0.6807
##    P-Value [Acc > NIR] : 0.04490
##
##          Kappa : 0.3928
##
## Mcnemar's Test P-Value : 0.06332
##
##          Sensitivity : 0.8889
##          Specificity : 0.4737
##    Pos Pred Value : 0.7826
##    Neg Pred Value : 0.6667
##          Prevalence : 0.6807
##    Detection Rate : 0.6050
##    Detection Prevalence : 0.7731
##    Balanced Accuracy : 0.6813
##
##    'Positive' Class : 0
##
```

*#using Decision Tree (rpart) model is giving 75.63% accuracy*

**### Build a random forest**

```
rf_model <- randomForest(admit ~ ., data=Train, proximity=FALSE,
                          ntree=15, mtry=3, na.action=na.omit)
rf_model
```

```
##
## Call:
## randomForest(formula = admit ~ ., data = Train, proximity = FALSE,      ntree = 15, mtry = 3, na.action = na.omit)
##          Type of random forest: classification
##          Number of trees: 15
## No. of variables tried at each split: 3
##
##          OOB estimate of  error rate: 36.96%
## Confusion matrix:
##      0  1 class.error
## 0 141 47      0.250
## 1  55 33      0.625
```

```
summary(rf_model)
```

```
##          Length Class  Mode
## call          7      -none- call
## type           1      -none- character
## predicted     276     factor numeric
## err.rate       45     -none- numeric
## confusion       6     -none- numeric
## votes         552     matrix numeric
```

```
## oob.times      276    -none- numeric
## classes        2    -none- character
## importance      3    -none- numeric
## importanceSD    0    -none- NULL
## localImportance 0    -none- NULL
## proximity       0    -none- NULL
## ntree           1    -none- numeric
## mtry            1    -none- numeric
## forest          14    -none- list
## y               276    factor numeric
## test            0    -none- NULL
## inbag           0    -none- NULL
## terms           3     terms  call
```

```
#using the model to predict on test data and using confusion matrix to calculate accuracy
rf_pred <- predict(rf_model, newdata=Test[, -1], type='Class')
```

```
confusionMatrix(rf_pred, Test$admit, positive = "0")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 70 27
##           1 11 11
##
##           Accuracy : 0.6807
##           95% CI : (0.589, 0.7631)
##           No Information Rate : 0.6807
##           P-Value [Acc > NIR] : 0.54380
##
##           Kappa : 0.173
##
##           Mcnemar's Test P-Value : 0.01496
##
##           Sensitivity : 0.8642
##           Specificity : 0.2895
##           Pos Pred Value : 0.7216
##           Neg Pred Value : 0.5000
##           Prevalence : 0.6807
##           Detection Rate : 0.5882
##           Detection Prevalence : 0.8151
##           Balanced Accuracy : 0.5768
##
##           'Positive' Class : 0
##
```

```
#using Random Forrest model is giving 68.07% accuracy
```

```
# build the Support Vector Machines(SVM) model
svm_model<- ksvm(admit ~ ., data = Train, scale = FALSE , C=25)
summary(svm_model)
```

```
## Length Class Mode
```



```
##      1   ksvm   S4
```

```
# Predicting the model results
```

```
svm_predict <- predict(svm_model, Test[, -1])
```

```
#SVM model accuracy using confusion matrix
```

```
confusionMatrix(svm_predict, Test$admit)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0  1
```

```
##           0 75 30
```

```
##           1  6  8
```

```
##
```

```
##           Accuracy : 0.6975
```

```
##           95% CI : (0.6065, 0.7783)
```

```
## No Information Rate : 0.6807
```

```
## P-Value [Acc > NIR] : 0.3882638
```

```
##
```

```
##           Kappa : 0.1639
```

```
##
```

```
## McNemar's Test P-Value : 0.0001264
```

```
##
```

```
##           Sensitivity : 0.9259
```

```
##           Specificity : 0.2105
```

```
## Pos Pred Value : 0.7143
```

```
## Neg Pred Value : 0.5714
```

```
## Prevalence : 0.6807
```

```
## Detection Rate : 0.6303
```

```
## Detection Prevalence : 0.8824
```

```
## Balanced Accuracy : 0.5682
```

```
##
```

```
## 'Positive' Class : 0
```

```
##
```

```
#using SVM model is giving 69.75% accuracy
```

```
# 10. Determine the accuracy rates for each kind of model
```

```
# a) Logistic Regression : 73.11%
```

```
# b) Decision Tree (C50) : 68.07%
```

```
# c) Decision Tree (rpart) : 75.63%
```

```
# d) Random Forest : 68.07%
```

```
# e) SVM : 69.75%
```

```
# 11. Select the most accurate model
```

```
# Decision - Decision Tree using library rpart is giving best accuracy
```

```
# which is 75.63% compared to other models
```

```
# So Most accurate model is Decision Tree(Rpart) Model
```

```
# 12. Identify other Machine learning or statistical techniques
# I have already used Random Forest. Other than random forest,
# we can apply Naive Bayes and Boosting Algorithms.
```

```
# Naive Bayes
library(naivebayes)
```

```
## naivebayes 0.9.7 loaded
```

```
naive_bayes_model <- naive_bayes(admit~.,Train)
naive_bayes_model
```

```
##
## ===== Naive Bayes =====
##
## Call:
## naive_bayes.formula(formula = admit ~ ., data = Train)
##
## -----
##
## Laplace smoothing: 0
##
## -----
##
## A priori probabilities:
##
##      0      1
## 0.6811594 0.3188406
##
## -----
##
## Tables:
##
## -----
## ::: gre (Gaussian)
## -----
##
## gre      0      1
## mean -0.07402582 0.23935859
## sd    1.03548586 0.97246981
##
## -----
## ::: gpa (Gaussian)
## -----
##
## gpa      0      1
## mean -0.0538757 0.1880284
## sd    1.0229455 1.0078305
##
## -----
## ::: rank (Categorical)
## -----
```

```
##
## rank      0      1
##   1 0.11702128 0.20454545
##   2 0.36702128 0.48863636
##   3 0.32446809 0.21590909
##   4 0.19148936 0.09090909
##
## -----
```

```
pred_nb <- predict(naive_bayes_model,Test[,-1])
confusionMatrix(pred_nb,Test$admit)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 76 28
##           1  5 10
##
##           Accuracy : 0.7227
##           95% CI : (0.6332, 0.8008)
##       No Information Rate : 0.6807
##       P-Value [Acc > NIR] : 0.1888403
##
##           Kappa : 0.24
##
##  Mcnemar's Test P-Value : 0.0001283
##
##           Sensitivity : 0.9383
##           Specificity : 0.2632
##           Pos Pred Value : 0.7308
##           Neg Pred Value : 0.6667
##           Prevalence : 0.6807
##           Detection Rate : 0.6387
##       Detection Prevalence : 0.8739
##           Balanced Accuracy : 0.6007
##
##           'Positive' Class : 0
##
```

```
train.control <- trainControl(method = "repeatedcv", number = 5, repeats = 3)
model_xgbTree <- train(admit ~ .,data=Train, method = "xgbTree",
                      trControl = train.control, verbosity = 0)
print(model_xgbTree)
```

```
## eXtreme Gradient Boosting
##
## 276 samples
## 3 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 3 times)
```

## Summary of sample sizes: 222, 221, 220, 221, 220, 221, ...

## Resampling results across tuning parameters:

##

##	eta	max_depth	colsample_bytree	subsample	nrounds	Accuracy
##	0.3	1	0.6	0.50	50	0.6508842
##	0.3	1	0.6	0.50	100	0.6533750
##	0.3	1	0.6	0.50	150	0.6558907
##	0.3	1	0.6	0.75	50	0.6545863
##	0.3	1	0.6	0.75	100	0.6593931
##	0.3	1	0.6	0.75	150	0.6543899
##	0.3	1	0.6	1.00	50	0.6377481
##	0.3	1	0.6	1.00	100	0.6520314
##	0.3	1	0.6	1.00	150	0.6533101
##	0.3	1	0.8	0.50	50	0.6667797
##	0.3	1	0.8	0.50	100	0.6435915
##	0.3	1	0.8	0.50	150	0.6579181
##	0.3	1	0.8	0.75	50	0.6460574
##	0.3	1	0.8	0.75	100	0.6520539
##	0.3	1	0.8	0.75	150	0.6507095
##	0.3	1	0.8	1.00	50	0.6498060
##	0.3	1	0.8	1.00	100	0.6472703
##	0.3	1	0.8	1.00	150	0.6520314
##	0.3	2	0.6	0.50	50	0.6520771
##	0.3	2	0.6	0.50	100	0.6557359
##	0.3	2	0.6	0.50	150	0.6435490
##	0.3	2	0.6	0.75	50	0.6484159
##	0.3	2	0.6	0.75	100	0.6218350
##	0.3	2	0.6	0.75	150	0.6302998
##	0.3	2	0.6	1.00	50	0.6386957
##	0.3	2	0.6	1.00	100	0.6327209
##	0.3	2	0.6	1.00	150	0.6338015
##	0.3	2	0.8	0.50	50	0.6437903
##	0.3	2	0.8	0.50	100	0.6304970
##	0.3	2	0.8	0.50	150	0.6449118
##	0.3	2	0.8	0.75	50	0.6399319
##	0.3	2	0.8	0.75	100	0.6449118
##	0.3	2	0.8	0.75	150	0.6399976
##	0.3	2	0.8	1.00	50	0.6241711
##	0.3	2	0.8	1.00	100	0.6228467
##	0.3	2	0.8	1.00	150	0.6240580
##	0.3	3	0.6	0.50	50	0.6330319
##	0.3	3	0.6	0.50	100	0.6485714
##	0.3	3	0.6	0.50	150	0.6305820
##	0.3	3	0.6	0.75	50	0.6447379
##	0.3	3	0.6	0.75	100	0.6544813
##	0.3	3	0.6	0.75	150	0.6411456
##	0.3	3	0.6	1.00	50	0.6448244
##	0.3	3	0.6	1.00	100	0.6315095
##	0.3	3	0.6	1.00	150	0.6350802
##	0.3	3	0.8	0.50	50	0.6498244
##	0.3	3	0.8	0.50	100	0.6534400
##	0.3	3	0.8	0.50	150	0.6364895
##	0.3	3	0.8	0.75	50	0.6243250
##	0.3	3	0.8	0.75	100	0.6377056

##	0.3	3	0.8	0.75	150	0.6400425
##	0.3	3	0.8	1.00	50	0.6204209
##	0.3	3	0.8	1.00	100	0.6242360
##	0.3	3	0.8	1.00	150	0.6182820
##	0.4	1	0.6	0.50	50	0.6559548
##	0.4	1	0.6	0.50	100	0.6557568
##	0.4	1	0.6	0.50	150	0.6303399
##	0.4	1	0.6	0.75	50	0.6497379
##	0.4	1	0.6	0.75	100	0.6582027
##	0.4	1	0.6	0.75	150	0.6556918
##	0.4	1	0.6	1.00	50	0.6438087
##	0.4	1	0.6	1.00	100	0.6520539
##	0.4	1	0.6	1.00	150	0.6496296
##	0.4	1	0.8	0.50	50	0.6472936
##	0.4	1	0.8	0.50	100	0.6582027
##	0.4	1	0.8	0.50	150	0.6486813
##	0.4	1	0.8	0.75	50	0.6520531
##	0.4	1	0.8	0.75	100	0.6435458
##	0.4	1	0.8	0.75	150	0.6543699
##	0.4	1	0.8	1.00	50	0.6461889
##	0.4	1	0.8	1.00	100	0.6520539
##	0.4	1	0.8	1.00	150	0.6556253
##	0.4	2	0.6	0.50	50	0.6388520
##	0.4	2	0.6	0.50	100	0.6342007
##	0.4	2	0.6	0.50	150	0.6316450
##	0.4	2	0.6	0.75	50	0.6459492
##	0.4	2	0.6	0.75	100	0.6396697
##	0.4	2	0.6	0.75	150	0.6412570
##	0.4	2	0.6	1.00	50	0.6277625
##	0.4	2	0.6	1.00	100	0.6324587
##	0.4	2	0.6	1.00	150	0.6374619
##	0.4	2	0.8	0.50	50	0.6532668
##	0.4	2	0.8	0.50	100	0.6364510
##	0.4	2	0.8	0.50	150	0.6350184
##	0.4	2	0.8	0.75	50	0.6302301
##	0.4	2	0.8	0.75	100	0.6169857
##	0.4	2	0.8	0.75	150	0.6252734
##	0.4	2	0.8	1.00	50	0.6350794
##	0.4	2	0.8	1.00	100	0.6159027
##	0.4	2	0.8	1.00	150	0.6169160
##	0.4	3	0.6	0.50	50	0.6498068
##	0.4	3	0.6	0.50	100	0.6461905
##	0.4	3	0.6	0.50	150	0.6317981
##	0.4	3	0.6	0.75	50	0.6351908
##	0.4	3	0.6	0.75	100	0.6364727
##	0.4	3	0.6	0.75	150	0.6412530
##	0.4	3	0.6	1.00	50	0.6266178
##	0.4	3	0.6	1.00	100	0.6291077
##	0.4	3	0.6	1.00	150	0.6327882
##	0.4	3	0.8	0.50	50	0.6376583
##	0.4	3	0.8	0.50	100	0.6447603
##	0.4	3	0.8	0.50	150	0.6292200
##	0.4	3	0.8	0.75	50	0.6486155
##	0.4	3	0.8	0.75	100	0.6354089

##	0.4	3	0.8	0.75	150	0.6366651
##	0.4	3	0.8	1.00	50	0.6172046
##	0.4	3	0.8	1.00	100	0.6232660
##	0.4	3	0.8	1.00	150	0.6220747
##	Kappa					
##	3.685579e-02					
##	7.534871e-02					
##	1.216743e-01					
##	2.627343e-02					
##	7.990407e-02					
##	8.780889e-02					
##	-3.917632e-02					
##	3.860905e-02					
##	5.118630e-02					
##	6.364288e-02					
##	5.952866e-02					
##	1.084836e-01					
##	1.002143e-02					
##	6.068873e-02					
##	9.130938e-02					
##	-9.634290e-03					
##	1.828988e-02					
##	4.189686e-02					
##	1.019261e-01					
##	1.333316e-01					
##	1.227399e-01					
##	9.345025e-02					
##	5.982175e-02					
##	9.212853e-02					
##	6.529361e-02					
##	7.392589e-02					
##	9.559793e-02					
##	5.799685e-02					
##	7.372039e-02					
##	1.193502e-01					
##	5.170855e-02					
##	1.080348e-01					
##	1.065808e-01					
##	1.542224e-02					
##	5.630381e-02					
##	7.500519e-02					
##	7.964340e-02					
##	1.358438e-01					
##	1.036282e-01					
##	1.121338e-01					
##	1.511013e-01					
##	1.404052e-01					
##	9.523142e-02					
##	9.093560e-02					
##	1.074978e-01					
##	1.270032e-01					
##	1.486643e-01					
##	1.181953e-01					
##	5.508229e-02					

## 1.158355e-01  
## 1.361926e-01  
## 3.660436e-02  
## 7.677235e-02  
## 7.197338e-02  
## 8.040067e-02  
## 1.178564e-01  
## 5.526494e-02  
## 4.298114e-02  
## 8.940785e-02  
## 1.123139e-01  
## 4.762013e-04  
## 5.563008e-02  
## 5.862813e-02  
## 3.689814e-02  
## 9.612135e-02  
## 1.014035e-01  
## 5.354440e-02  
## 7.950124e-02  
## 1.293429e-01  
## 3.069677e-05  
## 4.192773e-02  
## 6.904182e-02  
## 8.881188e-02  
## 9.973840e-02  
## 1.137594e-01  
## 9.052183e-02  
## 1.097942e-01  
## 1.229047e-01  
## 4.372097e-02  
## 8.909687e-02  
## 1.141158e-01  
## 1.191130e-01  
## 1.086639e-01  
## 1.151491e-01  
## 5.801564e-02  
## 7.065358e-02  
## 1.007069e-01  
## 6.411381e-02  
## 4.156881e-02  
## 5.861449e-02  
## 1.339567e-01  
## 1.312748e-01  
## 1.065693e-01  
## 9.716368e-02  
## 1.247375e-01  
## 1.326003e-01  
## 7.230958e-02  
## 9.559686e-02  
## 1.165391e-01  
## 8.226904e-02  
## 1.309609e-01  
## 1.050572e-01  
## 1.229370e-01

```
## 1.080579e-01
## 1.297676e-01
## 6.259913e-02
## 8.575536e-02
## 8.882034e-02
##
## Tuning parameter 'gamma' was held constant at a value of 0
## Tuning
## parameter 'min_child_weight' was held constant at a value of 1
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were nrounds = 50, max_depth = 1, eta
## = 0.3, gamma = 0, colsample_bytree = 0.8, min_child_weight = 1 and subsample
## = 0.5.
```

```
pred_xgboost <- predict(model_xgbTree,Test[,-1])
confusionMatrix(pred_xgboost,Test$admit)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 77 29
##           1  4  9
##
##           Accuracy : 0.7227
##           95% CI : (0.6332, 0.8008)
##           No Information Rate : 0.6807
##           P-Value [Acc > NIR] : 0.1888
##
##           Kappa : 0.2271
##
## Mcnemar's Test P-Value : 2.943e-05
##
##           Sensitivity : 0.9506
##           Specificity : 0.2368
##           Pos Pred Value : 0.7264
##           Neg Pred Value : 0.6923
##           Prevalence : 0.6807
##           Detection Rate : 0.6471
##           Detection Prevalence : 0.8908
##           Balanced Accuracy : 0.5937
##
##           'Positive' Class : 0
##
```

```
# Descriptive:
# Categorize the average of grade point into High, Medium,
#and Low (with admission probability percentages) and plot it on a point chart.
# Cross grid for admission variables with GRE Categorization is shown below:
```

```
# GRE Categorized 0-440 Low 440-580 Medium 580+ High
```

```
head(clgad)
```



```
##   admit gre  gpa ses Gender_Male Race rank
## 1     0 380 3.61  1           0    3    3
## 2     1 660 3.67  2           0    2    3
## 3     1 800 4.00  2           0    2    1
## 4     1 640 3.19  1           1    2    4
## 5     0 520 2.93  3           1    2    4
## 6     1 760 3.00  2           1    1    2
```

```
clgad$Category[clgad$gre < 440] <- "Low"
clgad$Category[clgad$gre >= 440 & clgad$gre < 580] <- "Medium"
clgad$Category[clgad$gre >= 580] <- "High"
head(clgad)
```

```
##   admit gre  gpa ses Gender_Male Race rank Category
## 1     0 380 3.61  1           0    3    3      Low
## 2     1 660 3.67  2           0    2    3      High
## 3     1 800 4.00  2           0    2    1      High
## 4     1 640 3.19  1           1    2    4      High
## 5     0 520 2.93  3           1    2    4    Medium
## 6     1 760 3.00  2           1    1    2      High
```

```
clgad$Category <- as.factor(clgad$Category)
str(clgad)
```

```
## 'data.frame':   395 obs. of  8 variables:
## $ admit      : Factor w/ 2 levels "0","1": 1 2 2 2 1 2 2 1 2 1 ...
## $ gre        : int  380 660 800 640 520 760 560 400 540 700 ...
## $ gpa        : num  3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
## $ ses        : Factor w/ 3 levels "1","2","3": 1 2 2 1 3 2 2 2 1 1 ...
## $ Gender_Male: Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 2 1 ...
## $ Race       : Factor w/ 3 levels "1","2","3": 3 2 2 2 2 1 2 2 1 2 ...
## $ rank       : Factor w/ 4 levels "1","2","3","4": 3 3 1 4 4 2 1 2 3 2 ...
## $ Category   : Factor w/ 3 levels "High","Low","Medium": 2 1 1 1 3 1 3 2 3 1 ...
```

```
summary(clgad$Category)
```

```
##   High    Low Medium
##   226     33   136
```

```
plot(clgad$Category)
```

